

API Gateway Implementation Steps - FreedmAI CI/CD System

Overview

Prerequisites

Phase 1: Repository Setup

- 1.1 Create GitHub Repository
- 1.2 Copy Implementation Files
- 1.3 Initial Commit

Phase 2: AWS Infrastructure Setup

- 2.1 Deploy Terraform Infrastructure
- 2.2 Note Important Outputs

Phase 3: GitHub Configuration

- 3.1 Configure Repository Secrets
- 3.2 Configure Environments
- 3.3 Enable GitHub Actions

Phase 4: EC2 Instance Setup

- 4.1 Launch EC2 Instances
- 4.2 Configure EC2 Instances
- 4.3 Configure IAM Role for EC2

Phase 5: Load Balancer Setup

- 5.1 Create Application Load Balancer
- 5.2 Configure Listeners

Phase 6: Nginx Configuration

- 6.1 Create Nginx Configuration Files

Phase 7: First Deployment

- 7.1 Trigger CI Pipeline
- 7.2 Deploy to UAT
- 7.3 Verify UAT Deployment

Phase 8: Production Deployment

- 8.1 UAT Sign-off Process
- 8.2 Deploy to Production
- 8.3 Verify Production Deployment

Phase 9: Monitoring Setup

- 9.1 CloudWatch Configuration
- 9.2 Alerting Setup

Phase 10: Security Hardening

- 10.1 Security Group Configuration
- 10.2 WAF Configuration

Phase 11: Cost Optimization

- 11.1 Resource Monitoring
- 11.2 Lifecycle Policies

Verification Checklist

- ☐ Infrastructure
- ☐ CI/CD Pipeline
- ☐ Deployment
- ☐ Monitoring
- ☐ Security

Troubleshooting Guide

Common Issues

Next Steps

API Gateway Implementation Steps - FreedmAI CI/CD System

Overview

This document provides step-by-step instructions for implementing the API Gateway service as the first microservice in the FreedmAI CI/CD system. The implementation includes complete CI/CD pipeline, security scanning, and deployment automation.

Prerequisites

- AWS Account with appropriate permissions
- GitHub organization: freedmai
- Local development environment with Node.js, Docker, and Terraform
- AWS CLI configured
- GitHub CLI installed

Phase 1: Repository Setup

1.1 Create GitHub Repository

```
# Create repository in freedmai organization
gh repo create freedmai/api-gateway --public --description "API
Gateway for FreedmAI microservices"

# Clone the repository
git clone https://github.com/freedmai/api-gateway.git
cd api-gateway
```

1.2 Copy Implementation Files

```
# Copy all files from local implementation
cp -r /var/Freedm/project/api-gateway/* .
cp -r /var/Freedm/project/api-gateway/. * . 2>/dev/null || true

# Remove node_modules if copied
rm -rf node_modules/
```

1.3 Initial Commit

```
git add .
git commit -m "Initial API Gateway implementation with CI/CD
pipeline"
git push origin main
```

Phase 2: AWS Infrastructure Setup

2.1 Deploy Terraform Infrastructure

```
cd terraform/

# Initialize Terraform
terraform init

# Review the plan
terraform plan

# Apply infrastructure
terraform apply
```

Expected Resources Created: - ECR Repository: freedmai-api-gateway -
IAM Role: GitHubActionsRole-FreedmAI - OIDC Provider for GitHub Actions

2.2 Note Important Outputs

```
# Save these values for GitHub secrets
terraform output ecr_repository_url
terraform output github_actions_role_arn
```

Phase 3: GitHub Configuration

3.1 Configure Repository Secrets

```
# Add AWS role ARN
gh secret set AWS_ROLE_ARN --body
"arn:aws:iam::ACCOUNT_ID:role/GitHubActionsRole-FreedmAI"

# Add ECR registry URL
gh secret set ECR_REGISTRY --body "ACCOUNT_ID.dkr.ecr.us-east-
1.amazonaws.com"
```

3.2 Configure Environments

```
# Create UAT environment
gh api repos/freedmai/api-gateway/environments/uat --method PUT

# Create Production environment with protection rules
gh api repos/freedmai/api-gateway/environments/production --method
PUT \
  --field
  protection_rules='[{"type":"required_reviewers","reviewers":
  [{"type":"Team","id":TEAM_ID}]}'
```

3.3 Enable GitHub Actions

```
# Enable Actions if not already enabled
gh api repos/freedmai/api-gateway/actions/permissions --method PUT \
  --field enabled=true --field allowed_actions=all
```

Phase 4: EC2 Instance Setup

4.1 Launch EC2 Instances

UAT Instance (t3.small):

```
# Launch UAT instance
aws ec2 run-instances \
  --image-id ami-0c02fb55956c7d316 \
  --instance-type t3.small \
  --key-name your-key-pair \
  --security-group-ids sg-xxxxxxx \
  --subnet-id subnet-xxxxxxx \
  --tag-specifications 'ResourceType=instance,Tags=
[ {Key=Name,Value=FreedmAI-UAT}, {Key=Environment,Value=uat} ] '
```

Production Instance (t3.small):

```
# Launch Production instance
aws ec2 run-instances \
  --image-id ami-0c02fb55956c7d316 \
  --instance-type t3.small \
  --key-name your-key-pair \
  --security-group-ids sg-xxxxxxx \
  --subnet-id subnet-xxxxxxx \
  --tag-specifications 'ResourceType=instance,Tags=
[ {Key=Name,Value=FreedmAI-Prod}, {Key=Environment,Value=production} ] '
```

4.2 Configure EC2 Instances

For both UAT and Production instances:

```
# SSH into instance
ssh -i your-key.pem ec2-user@INSTANCE_IP

# Update system
sudo yum update -y

# Install Docker
sudo yum install -y docker
sudo systemctl start docker
sudo systemctl enable docker
sudo usermod -a -G docker ec2-user

# Install Docker Compose
sudo curl -L
"https://github.com/docker/compose/releases/latest/download/docker-
compose-$(uname -s)-$(uname -m)" -o /usr/local/bin/docker-compose
sudo chmod +x /usr/local/bin/docker-compose

# Install AWS CLI v2
curl "https://awscli.amazonaws.com/awscli-exe-linux-x86_64.zip" -o
"awscliv2.zip"
unzip awscliv2.zip
sudo ./aws/install

# Install GitHub Actions Runner (UAT only)
mkdir actions-runner && cd actions-runner
```

```

curl -o actions-runner-linux-x64-2.311.0.tar.gz -L
https://github.com/actions/runner/releases/download/v2.311.0/actions-
runner-linux-x64-2.311.0.tar.gz
tar xzf ./actions-runner-linux-x64-2.311.0.tar.gz

# Configure runner (get token from GitHub)
./config.sh --url https://github.com/freedmai/api-gateway --token
YOUR_TOKEN

# Install as service
sudo ./svc.sh install
sudo ./svc.sh start

```

4.3 Configure IAM Role for EC2

```

# Create IAM role for EC2 instances
aws iam create-role --role-name FreedmAI-EC2-Role --assume-role-
policy-document '{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": {
        "Service": "ec2.amazonaws.com"
      },
      "Action": "sts:AssumeRole"
    }
  ]
}'

# Attach ECR permissions
aws iam attach-role-policy --role-name FreedmAI-EC2-Role --policy-
arn arn:aws:iam::aws:policy/AmazonEC2ContainerRegistryReadOnly

# Create instance profile
aws iam create-instance-profile --instance-profile-name FreedmAI-
EC2-Profile
aws iam add-role-to-instance-profile --instance-profile-name
FreedmAI-EC2-Profile --role-name FreedmAI-EC2-Role

# Attach to instances
aws ec2 associate-iam-instance-profile --instance-id i-xxxxxxx --
iam-instance-profile Name=FreedmAI-EC2-Profile

```

Phase 5: Load Balancer Setup

5.1 Create Application Load Balancer

```

# Create ALB
aws elbv2 create-load-balancer \
  --name FreedmAI-ALB \
  --subnets subnet-xxxxxxx subnet-yyyyyyy \
  --security-groups sg-xxxxxxx \
  --scheme internet-facing \
  --type application \
  --ip-address-type ipv4

```

```

# Create target groups
aws elbv2 create-target-group \
  --name FreedmAI-UAT-TG \
  --protocol HTTP \
  --port 80 \
  --vpc-id vpc-xxxxxxxx \
  --health-check-path /health

aws elbv2 create-target-group \
  --name FreedmAI-Prod-TG \
  --protocol HTTP \
  --port 80 \
  --vpc-id vpc-xxxxxxxx \
  --health-check-path /health

# Register targets
aws elbv2 register-targets --target-group-arn
arn:aws:elasticloadbalancing:us-east-1:ACCOUNT:targetgroup/FreedmAI-
UAT-TG/xxxxxxxx --targets Id=i-xxxxxxxx
aws elbv2 register-targets --target-group-arn
arn:aws:elasticloadbalancing:us-east-1:ACCOUNT:targetgroup/FreedmAI-
Prod-TG/xxxxxxxx --targets Id=i-yyyyyyyyy

```

5.2 Configure Listeners

```

# Create listeners with path-based routing
aws elbv2 create-listener \
  --load-balancer-arn arn:aws:elasticloadbalancing:us-east-
1:ACCOUNT:loadbalancer/app/FreedmAI-ALB/xxxxxxxx \
  --protocol HTTP \
  --port 80 \
  --default-actions
Type=forward,TargetGroupArn=arn:aws:elasticloadbalancing:us-east-
1:ACCOUNT:targetgroup/FreedmAI-Prod-TG/xxxxxxxx

# Add rules for UAT routing
aws elbv2 create-rule \
  --listener-arn arn:aws:elasticloadbalancing:us-east-
1:ACCOUNT:listener/app/FreedmAI-ALB/xxxxxxxx/yyyyyyyy \
  --priority 100 \
  --conditions Field=host-header,Values=uat.freedmai.com \
  --actions
Type=forward,TargetGroupArn=arn:aws:elasticloadbalancing:us-east-
1:ACCOUNT:targetgroup/FreedmAI-UAT-TG/xxxxxxxx

```

Phase 6: Nginx Configuration

6.1 Create Nginx Configuration Files

On UAT Instance:

```

sudo mkdir -p /opt/freedmai/nginx
sudo tee /opt/freedmai/nginx/uat.conf > /dev/null <<EOF
events {
    worker_connections 1024;
}

```

```

http {
    upstream api_gateway {
        server api-gateway:3000;
    }

    server {
        listen 80;
        server_name uat.freedmai.com;

        location /health {
            proxy_pass http://api_gateway/health;
            proxy_set_header Host \ $host;
            proxy_set_header X-Real-IP \ $remote_addr;
        }

        location / {
            proxy_pass http://api_gateway;
            proxy_set_header Host \ $host;
            proxy_set_header X-Real-IP \ $remote_addr;
            proxy_set_header X-Forwarded-For
\ $proxy_add_x_forwarded_for;
            proxy_set_header X-Forwarded-Proto \ $scheme;
        }
    }
}
EOF

```

On Production Instance:

```

sudo mkdir -p /opt/freedmai/nginx
sudo tee /opt/freedmai/nginx/prod.conf > /dev/null <<EOF
events {
    worker_connections 1024;
}

http {
    upstream api_gateway {
        server api-gateway:3000;
    }

    server {
        listen 80;
        server_name api.freedmai.com;

        location /health {
            proxy_pass http://api_gateway/health;
            proxy_set_header Host \ $host;
            proxy_set_header X-Real-IP \ $remote_addr;
        }

        location / {
            proxy_pass http://api_gateway;
            proxy_set_header Host \ $host;
            proxy_set_header X-Real-IP \ $remote_addr;
            proxy_set_header X-Forwarded-For
\ $proxy_add_x_forwarded_for;
            proxy_set_header X-Forwarded-Proto \ $scheme;
        }
    }
}

```

EOF

Phase 7: First Deployment

7.1 Trigger CI Pipeline

```
# Push code to trigger CI
git add .
git commit -m "Configure deployment settings"
git push origin main
```

Expected CI Pipeline Steps: 1. □ Code checkout 2. □ Node.js setup and dependency installation 3. □ ESLint code quality check 4. □ Jest unit tests 5. □ npm security audit 6. □ Trivy filesystem vulnerability scan 7. □ GitLeaks secret scan 8. □ Docker image build and push to ECR 9. □ Docker image vulnerability scan

7.2 Deploy to UAT

```
# Trigger UAT deployment via GitHub Actions
gh workflow run deploy-uat.yml -f image_tag=latest
```

Expected UAT Deployment Steps: 1. □ AWS credentials configuration 2. □ ECR login 3. □ Docker image pull 4. □ Container deployment with Docker Compose 5. □ Health check verification 6. □ Deployment notification

7.3 Verify UAT Deployment

```
# Check container status
ssh ec2-user@UAT_INSTANCE_IP "docker ps"

# Test health endpoint
curl http://uat.freedmai.com/health

# Check logs
ssh ec2-user@UAT_INSTANCE_IP "docker logs freedmai-api-gateway-uat"
```

Phase 8: Production Deployment

8.1 UAT Sign-off Process

1. **Manual Testing:** Verify all functionality in UAT
2. **Performance Testing:** Load test the UAT environment
3. **Security Validation:** Confirm security scans passed
4. **Stakeholder Approval:** Get approval from designated approvers

8.2 Deploy to Production

```
# Trigger production deployment (requires approval)
gh workflow run deploy-prod.yml -f image_tag=latest -f
uat_approval_id=12345
```


Expected Production Deployment Steps: 1. ☐ UAT approval validation 2. ☐ AWS credentials configuration 3. ☐ Current version backup 4. ☐ Blue-green deployment 5. ☐ Health check verification 6. ☐ Traffic switch 7. ☐ Old container cleanup 8. ☐ Rollback capability on failure

8.3 Verify Production Deployment

```
# Test production endpoint
curl http://api.freedmai.com/health

# Monitor logs
ssh ec2-user@PROD_INSTANCE_IP "docker logs freedmai-api-gateway-
prod"
```

Phase 9: Monitoring Setup

9.1 CloudWatch Configuration

```
# Create log groups
aws logs create-log-group --log-group-name /freedmai/api-gateway/uat
aws logs create-log-group --log-group-name /freedmai/api-
gateway/production

# Create custom metrics
aws cloudwatch put-metric-data \
  --namespace "FreedmAI/APIGateway" \
  --metric-data MetricName=HealthCheck,Value=1,Unit=Count
```

9.2 Alerting Setup

```
# Create SNS topic
aws sns create-topic --name FreedmAI-Alerts

# Create CloudWatch alarms
aws cloudwatch put-metric-alarm \
  --alarm-name "API-Gateway-High-Error-Rate" \
  --alarm-description "API Gateway error rate too high" \
  --metric-name ErrorRate \
  --namespace FreedmAI/APIGateway \
  --statistic Average \
  --period 300 \
  --threshold 5.0 \
  --comparison-operator GreaterThanThreshold \
  --evaluation-periods 2
```

Phase 10: Security Hardening

10.1 Security Group Configuration

```
# Create security group for API Gateway
aws ec2 create-security-group \
  --group-name FreedmAI-API-Gateway-SG \
  --description "Security group for FreedmAI API Gateway"
```

```
# Allow HTTP traffic from ALB only
aws ec2 authorize-security-group-ingress \
  --group-id sg-xxxxxxx \
  --protocol tcp \
  --port 80 \
  --source-group sg-yyyyyyyyy
```

10.2 WAF Configuration

```
# Create WAF Web ACL
aws wafv2 create-web-acl \
  --name FreedmAI-WAF \
  --scope REGIONAL \
  --default-action Allow={} \
  --rules file://waf-rules.json
```

Phase 11: Cost Optimization

11.1 Resource Monitoring

```
# Set up cost alerts
aws budgets create-budget \
  --account-id ACCOUNT_ID \
  --budget file://budget-config.json
```

11.2 Lifecycle Policies

```
# ECR lifecycle policy
aws ecr put-lifecycle-policy \
  --repository-name freedmai-api-gateway \
  --lifecycle-policy-text file://ecr-lifecycle.json
```

Verification Checklist

☐ Infrastructure

- ☐ ECR repository created and accessible
- ☐ IAM roles and policies configured
- ☐ EC2 instances running and accessible
- ☐ Load balancer configured with health checks
- ☐ Security groups properly configured

☐ CI/CD Pipeline

- ☐ GitHub Actions workflows executing successfully
- ☐ Security scans passing (Trivy, GitLeaks, npm audit)
- ☐ Code quality checks passing (ESLint, tests)
- ☐ Docker images building and pushing to ECR
- ☐ Self-hosted runner operational on UAT instance

□ Deployment

- ☐ UAT deployment working via GitHub Actions
- ☐ Production deployment with approval gates
- ☐ Blue-green deployment strategy functional
- ☐ Rollback capability tested
- ☐ Health checks responding correctly

□ Monitoring

- ☐ CloudWatch logs streaming
- ☐ Custom metrics being recorded
- ☐ Alerts configured and tested
- ☐ Cost monitoring active

□ Security

- ☐ No secrets in code repository
- ☐ OIDC authentication working
- ☐ Container vulnerability scans passing
- ☐ Network security properly configured
- ☐ WAF rules active (if implemented)

Troubleshooting Guide

Common Issues

1. GitHub Actions failing with AWS permissions:

```
# Verify OIDC trust relationship
aws iam get-role --role-name GitHubActionsRole-FreedMAI
```

2. Docker container not starting:

```
# Check container logs
docker logs freedmai-api-gateway-uat
# Verify environment variables
docker exec freedmai-api-gateway-uat env
```

3. Health check failing:

```
# Test directly on instance
curl http://localhost:3000/health
# Check application logs
docker logs freedmai-api-gateway-uat
```

4. ECR push failing:

```
# Re-authenticate with ECR
aws ecr get-login-password --region us-east-1 | docker login --
username AWS --password-stdin ACCOUNT.dkr.ecr.us-east-
1.amazonaws.com
```

Next Steps

1. **Add More Microservices:** Replicate this pattern for other services
2. **Implement Deployment UI:** Create web interface for deployments
3. **Add Database Integration:** Set up RDS for shared configuration
4. **Enhance Monitoring:** Add application performance monitoring
5. **Implement Service Mesh:** Consider Istio or AWS App Mesh for advanced routing

Cost Summary

Monthly Estimated Costs: - EC2 instances (2x t3.small): ~\$30 - ALB: ~\$16 - ECR storage: ~\$2 - CloudWatch logs: ~\$5 - Data transfer: ~\$5 - **Total:** ~\$58/month

Support Contacts

- **DevOps Team:** devops@freedmai.com
- **Security Team:** security@freedmai.com
- **AWS Support:** Use AWS Support Center

Document Version: 1.0

Date: September 19, 2025

Implementation Status: Complete

Next Review: October 19, 2025